Liquid Democracy for Rating Systems

Hemanath Peddireddi

Department of Computer Science

University of Warwick

Supervised by Markus Brill

19 April 2025

WARWICK

THE UNIVERSITY OF WARWICK

Abstract

TODO: finish

This project explores how liquid democracy can be used to enhance rating systems by integrating it into vodle, an online polling platform where users rate options using sliders. Traditional liquid democracy models do not always reflect participants' true preferences, especially when some users abstain from voting or when some individuals gain too much influence. To address this, the project adds support for ranked delegation and weighted vote splitting.

Keywords:

Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1   Motivation - TODO (Finalise)

Liquid democracy has strong theoretical appeal as a flexible and participatory decision-making model, but practical implementations remain rare and underexplored. Most existing research assumes idealised conditions, while real-world systems must contend with asynchronous participation, limited user engagement, and the potential for structural issues like delegation cycles or vote loss.

1.2   Vodle - TODO: finalise

Vodle is an online platform where users participate in polls to vote on subjects through user created polls. Each poll contains a set of options, and users provide ratings for each option from 0 to 100, where a larger number means that they prefer the option more, using a slider. When the poll ends, the ratings submitted by voters are then aggregated and a result is calculated.

Vodle presents an opportunity to explore how liquid democracy can be adapted to work in practice. This project is motivated by the need to understand and resolve the challenges of integrating delegation-based voting into an existing system, while improving the platform's fairness, expressiveness, and overall user experience. **can insert screenshots**

1.3   Liquid Democracy

Liquid democracy is a decision-making system that combines elements of both direct and representative democracy that offers a voter more flexibility than traditional voting models.

In direct democracy, every participant votes individually on each issue. This model offers the most individual input but can become impractical for large-scale decision-making due to the high level of participation required from each individual. As Ford (2002) states, direct democracy assumes that all individuals are both willing and able to engage meaningfully with every decision, which is often not the case in large groups due to the variance in both the interest and knowledge of voters. The cognitive demand of staying informed on all matters, combined with the time commitment necessary for constant participation, makes direct democracy unmanageable at scale.

In a representative democracy, citizens elect officials who make decisions on their behalf for the duration of a fixed term. While this model is scalable and practical for large populations, it introduces several limitations. Elected representatives often make decisions based on party lines, personal convictions, or external influences such as lobbying groups, which may not accurately reflect the preferences of their constituents (Blum and Zuber, 2016). In addition, because elections are infrequent, this system tends to be unresponsive to shifts in public opinion. Citizens are unable to easily adjust or retract their delegation, which limits their ability to influence decisions once representatives are in office (Blum and Zuber, 2016). As a result, participation is both indirect and inflexible, which can lead to disengagement and dissatisfaction among voters.

Liquid democracy addresses these limitations by allowing voters to: cast their votes directly, delegate them to someone that they trust, or abstain from voting entirely (Blum and Zuber, 2016). In comparison to a direct democracy, the bar for participation is lowered as voters no longer need to stay informed and engaged to pass a vote because they can trust a delegate to do it on their behalf. These delegations can also be updated or revoked at any time, giving users more control over how their vote is used in comparison to a traditional representational democracy where your representative can only be changed at certain points in time.

1.4   Project Goal - TODO

The project's main goal is to integrate liquid democracy into the vodle platform.

Key features include ranked delegation and vote splitting . . .

1.5   Project Outline - TODO

This report is structured as follows: *will add when report is written*

Chapter 2

Background Research

This chapter provides background context for the development of a liquid democracy system within vodle. It builds on the concepts introduced earlier, focusing on more detailed research into known limitations of liquid democracy and potential solutions proposed in academic literature. Additionally, the technical foundations and design philosophy of vodle as a platform are explored.

Throughout this chapter, several diagrams are used to illustrate how votes move through a liquid democracy system. To clarify the roles of different voters within these diagrams, the following symbols are used:

- **Circle:** *Delegated voter* - a voter that has delegated and therefore does not cast their own vote.

- **Square:** *Casting voter* - a voter that has not delegated and casts their own vote.

- **Triangle:** *Abstaining voter* - a voter that neither delegates nor casts their own vote.

2.1   Liquid Democracy

Liquid democracy, or delegative voting, allows voters to either cast a vote directly, delegate it to someone they trust, or abstain (Blum and Zuber, 2016). A key feature is that delegations are transitive - a chain of users that all delegate to each other sequentially ends with a single final voter who casts their vote on behalf of all those in the chain.

Although transitivity provides greater flexibility to the voters, it can also lead to unintended consequences. Chains of delegations may result in cycles that prevent votes

from being cast, or allow certain individuals to accumulate an excessive share of influence, creating "super-voters". These problems amongst others motivate the need for alternative delegation mechanisms, as discussed later in this chapter.

2.1.1 Issues with Liquid Democracy

Delegation cycles

Figure 2.1: Delegation cycle: A delegates to B, B to C, and C back to A.

Delegation cycles occur when a vote is delegated in such a way that it ends up forming a loop (Brill et al., 2022), preventing the vote from reaching a final casting voter. For example, if Alice delegates her vote to Bob, Bob delegates to Charlie, and Charlie delegates back to Alice, the votes become trapped in a cycle (seen above) and can be treated as a loss of representation (Christoff and Grossi, 2017).

This issue is particularly problematic because it can nullify votes without the affected users ever realising. In systems where cycles are not explicitly detected and handled, these votes could be discarded silently, potentially changing the final outcome of the votes.

A simplistic method to prevent cycles is to check whether a delegation would create a cycle before allowing it. For example, if Alice tries to delegate her vote to Bob, the system checks whether Bob has already directly or indirectly delegated their vote to Alice. If so, the delegation is rejected. However, this approach can be cumbersome and may lead to a poor user experience, as users may not understand why their delegation was denied.

Delegation cycles are increasingly likely to emerge in dynamic voting systems, where delegations can be added, removed, or modified at any point in time. Delegations that initially did not form part of a cycle may later contribute to one as other voters add a new delegation or alter an existing one.

Abstentions

Figure 2.2: Delegation chain ending in abstention: A delegates to B, B to C. C abstains, causing the votes of A and B to be lost.

A voter abstains by neither casting a vote nor delegating it to another user (Brill et al., 2022). This includes both deliberate abstention, where a voter knowingly chooses not to participate, and passive abstention, where a voter may be unaware of an ongoing poll or are unable to engage with it.

Abstentions are especially impactful when they occur at the end of a larger delegation chain, as all votes passed along the chain to that voter are effectively lost (Brill et al., 2022). Additionally, the voters whose decisions were passed along the chain may also be unaware that their votes have been nullified, worsening the effect of the abstention.

Super-voters

Figure 2.3: Super-voter: A delegates to B, B to C. No matter which vote D or E cast, C's vote will always determine the outcome as it has a weight of 3.

In liquid democracy, a super-voter is an individual who receives a large number of delegated votes, therefore gaining disproportionate influence over decisions (Kling et al., 2015). While this behaviour may reflect voters' genuine preferences, it can lead to a concentration of power that goes against the intended egalitarianism and democratic ideals of liquid democracy.

Even if a particular implementation of liquid democracy allows users to alter their delegation at any time, in practice, many voters may not actively monitor or even know how their vote is being used. This can allow a small number of super-voters to dominate outcomes, especially in systems that do nothing to prevent large delegation chains.

Real-world examples of this phenomenon have been documented. In the German Pirate Party's use of LiquidFeedback, certain users received so many delegations that their votes were like "decrees" (Sven Becker, 2012; Kling et al., 2015) even though they were not elected officials. Despite Kling et al. (2015) noting that super-voters generally voted in line with the majority and therefore did not drastically affect the outcome

of the votes, the potential for individuals to single-handedly influence results still remains a concern.

This pattern is not limited to traditional online voting platforms. It can also be seen within decentralised autonomous organisations (DAOs) - blockchain-based entities where decisions are made collectively by token holders without central leadership. These organisations use token-based voting to decide on critical issues like protocol upgrades and funding allocations. Hall and Miyazaki (2024) studied 18 decentralised autonomous organisations (DAOs) and found that voting power was often concentrated in the hands of a few delegates. While most did not control a large share of all available tokens, low participation meant that their share of actual votes cast was disproportionately high. In several DAOs, the top five delegates accounted for over 50% of all votes cast, and in the DAO Gitcoin, this figure exceeded 90%.

2.1.2 Variations of Liquid Democracy

The challenges discussed in the previous section, such as delegation cycles, vote loss due to abstentions, and the emergence of super-voters, highlight inherent vulnerabilities in the standard liquid democracy model. To mitigate these issues, enhancements have been proposed that modify how delegations function. These include techniques that allow voters to specify multiple delegates or distribute their vote to multiple casting voters. Each approach introduces different trade-offs and requires algorithmic support to ensure sound and interpretable outcomes.

The following subsections present several such variations, along with the algorithms that can be used to implement them.

Ranked Delegation

Ranked delegation improves liquid democracy by allowing voters to list several trusted delegates in order of preference. Instead of choosing just one delegate, a voter can specify a ranked list so that if their top choice is unavailable (e.g. due to abstention or being a part of a delegation cycle) the system can use the next delegate specified.

Implementing ranked delegation requires a mechanism to decide among multiple possible delegation paths – a route that a vote can take through the delegation graph to reach a casting voter. This is done through a *delegation rule*, a function that, given a ranked delegation instance and a delegating voter, selects a unique path leading to a *casting voter* (Brill et al., 2022).

The following key properties help evaluate these delegation rules:

- **Guru Participation:** Ensures that a voter accepting delegated votes (a "guru") is never worse off by doing so. Receiving additional delegations should not decrease their influence over the final outcome (Kotsialou and Riley, 2020).

- **Confluence:** Guarantees that each delegating voter ends up with one clear and unambiguous delegation path. This property simplifies vote resolution and enhances transparency (Brill et al., 2022).

- **Copy Robustness:** Prevents strategic manipulation where a voter might mimic another's vote outside the system to gain extra influence. A copy-robust rule makes sure that duplicating a vote externally does not yield more combined power than a proper delegation (Brill et al., 2022; Behrens and Swierczek, 2015).

The literature considers several delegation rules, each with distinct trade-offs:

Depth-First Delegation (DFD): Selects the path beginning with the highest-ranked delegate, even if the resulting chain is long. Although it prioritizes individual trust preferences, DFD can violate guru participation (Kotsialou and Riley, 2020).

Breadth-First Delegation (BFD): Chooses the shortest available delegation path and uses rankings only to resolve ties. This approach usually produces direct, predictable chains and satisfies guru participation, although it might sometimes assign a vote to a lower-ranked delegate (Kotsialou and Riley, 2020; Brill et al., 2022).

MinSum: Balances path length and delegation quality by selecting the path with the lowest total sum of edge ranks. Due to this, MinSum avoids both unnecessarily long chains and poorly ranked delegations (Brill et al., 2022).

Diffusion: Constructs delegation paths in stages by assigning votes layer by layer based on the lowest available rank at each step. This method tends to avoid poor delegations but can sometimes produce unintuitive outcomes due to its tie-breaking procedure (Brill et al., 2022).

Leximax: Compares paths based on their worst-ranked edge. This ensures that especially low-ranked delegations are avoided early in the path while maintaining confluence (Brill et al., 2022).

BordaBranching: Takes a global view of the delegation graph by selecting a branching that minimizes the total rank across all delegation edges. It satisfies both guru participation and copy robustness, though it is more computationally intensive (Brill et al., 2022).

In summary, ranked delegation enhances liquid democracy by reducing the risk of lost votes. The choice of delegation rule not only affects system efficiency but also influences fairness and robustness. While simpler methods such as DFD and BFD are easier to implement, advanced rules like MinSum, Leximax, and BordaBranching

offer stronger guarantees and are better suited for practical deployment in platforms such as vodle.

For our implementation, MinSum will be chosen as the delegation rule because it offers a good trade-off between delegation quality, computational efficiency, and user interpretability. By selecting the path with the lowest total rank sum, MinSum prioritises higher ranked delegates while avoiding unnecessarily long or indirect delegation chains. This not only improves the quality of representation but also makes it clearer to users why a particular delegate was chosen as the path reflects their stated preferences in a straightforward way. Additionally, MinSum is more computationally efficient than alternatives like BordaBranching, making it a practical choice for deployment within the vodle platform.

Vote Splitting

Traditional delegation systems, which require voters to delegate their entire vote to a single individual, introduce significant risks such as vote loss through delegate abstentions, delegation cycles, and excessive concentration of voting power in the hands of a few super-voters. To mitigate these issues, vote splitting allows voters to distribute their voting power among multiple delegates. This approach provides greater flexibility and robustness while preserving voter intent more accurately.

Vote splitting offers several key advantages:

- **Increased resilience:** Distributing votes across multiple delegates reduces the impact of any single delegate abstaining or becoming unavailable, thus lowering the risk of vote loss.

- **Reduced concentration of power:** Allowing partial votes to different delegates decreases the likelihood of any single delegate becoming a super-voter.

- **Enhanced voter expression:** Voters can more precisely express their preferences and trust levels by allocating voting power proportionally to multiple individuals.

Several methodologies for implementing vote splitting have been explored in the literature, each with its strengths and weaknesses:

Equal Vote Distribution (Degrave, 2014)

Degrave's approach allows voters to distribute their votes evenly among multiple delegates (Degrave, 2014). Voters select a group of delegates, and their vote is equally

distributed amongst those that do not abstain. Although this system is intuitive and reduces the impact of abstentions, it lacks flexibility as voters cannot express differing trust levels towards each delegate. Additionally, a critical limitation is the inability for voters to allocate any portion of their vote to themselves, meaning voters are forced to either delegate their entire voting power or none of it, severely limiting personal control over the final vote.

Fractional Delegation (Bersetche et al., 2024)

Bersetche et al. introduce fractional delegation, allowing voters to explicitly assign different weights to each chosen delegate, including themselves (Bersetche, 2024). Delegates each receive a specified fraction of the voter's total voting power, reflecting the voter's nuanced trust and preference levels. This approach captures detailed voter preferences accurately and allows for greater personal agency compared to equal vote distribution. However, fractional delegation introduces additional complexity in managing and tracking these weighted delegations. Users must explicitly manage multiple numerical allocations, which may increase cognitive load and complicate user interfaces.

Trust Matrix Iterative Model

In the trust matrix model, voters define explicit trust values $\text{trust}_{i,j}$, representing how much voter i trusts voter j. Each voter also assigns their personal rating (self_i). The effective rating for each voter is computed iteratively using the equation:

$$\text{eff}_i = \text{trust}_{i,i} \cdot \text{self}_i + \sum_{j \neq i} \text{trust}_{i,j} \cdot \text{eff}_j$$

Here, each voter's trust values (including self-trust) must sum to at most 1. The iterative computation continues until the change in effective ratings between iterations falls below a predefined threshold ϵ. This approach offers the highest granularity and expressive power, allowing voters to precisely articulate nuanced trust relationships among multiple delegates. However, it comes with considerable computational complexity and potential convergence issues, especially in large networks with dense delegation relationships. Additionally, users may find it challenging to specify and manage such detailed trust matrices, negatively affecting usability.

Summary of Approaches

In summary, each vote-splitting method balances voter expressivity, computational complexity, user interface clarity, and resilience differently:

- **Equal Vote Distribution (Degrave)** excels in simplicity and ease of implementation, ensuring robustness through straightforward delegation. However, it significantly limits voter expression and prohibits voters from allocating votes to themselves.

- **Fractional Delegation (Bersetche)** provides greater flexibility, permitting detailed voter preference expression, including self-allocation of votes. This method increases both computational complexity and interface complexity.

- **Trust Matrix Iterative Model** offers the highest expressivity and detail in delegation relationships, capturing complex trust dynamics. However, this method entails substantial computational overhead and introduces complexity in terms of usability and understanding for voters.

Considering the priorities of the vodle platform–usability, intuitive interfaces, computational efficiency, and maintaining voter agency–the fractional delegation method described by Bersetche et al. is most suitable. It effectively balances flexibility and clarity, permitting detailed voter preferences without overwhelming the computational or cognitive capacities of the system and its users.

2.2   Existing Implementations of Liquid Democracy

To understand how liquid democracy can be integrated into vodle, it is important to examine how similar systems have been implemented in real-world contexts. This section explores two implementations, LiquidFeedback and Google Votes, that offer valuable insights into the technical, social, and usability challenges associated with applying liquid democracy at scale.

2.2.1   LiquidFeedback

LiquidFeedback is one of the earliest and most influential real-world implementations of liquid democracy. Developed as an open-source platform, it was notably adopted by the German Pirate Party in 2010 to facilitate internal policy-making through online participation (Behrens et al., 2014). The platform allowed members to submit proposals, debate them in structured phases, and vote either directly or via transitive delegation.

In LiquidFeedback, users could choose different delegates for different topics, allowing them to assign their vote to someone they trusted on a specific issue. These choices remained in place until the user changed them, which meant that certain individuals could gradually accumulate more influence if others did not update their delegations.

When multiple proposals were put forward, the system used a ranking-based voting method (such as the Schulze method) to decide which one should win. This approach compares each proposal against the others and selects the one that would win the most head-to-head matchups. Importantly, the system only accepted a proposal if it clearly beat the alternative of doing nothing, helping to avoid unnecessary or unpopular changes.

In practice, the Pirate Party's use of LiquidFeedback revealed several key dynamics relevant to this project. The platform was successful in enabling large-scale participation and crowdsourced policy formation, but it also demonstrated common risks of liquid democracy. Such as the existence of super-voters, as discussed previously.

Another practical issue was the complexity of the system. LiquidFeedback was difficult to understand for many users, especially those unfamiliar with concepts like transitive delegation or multi-stage voting which limited its accessibility and contributed to declining engagement over time (Kling et al., 2015).

For a platform like vodle, the experience of LiquidFeedback highlights several important design considerations. First, user interfaces must be intuitive enough to allow voters to participate without needing deep technical knowledge. Second, the user must be know the status of their delegation at a glance - improving the understanding of the platform. Finally, ensuring that votes lead to visible and actionable outcomes is critical for maintaining user engagement.

2.2.2 Google Votes
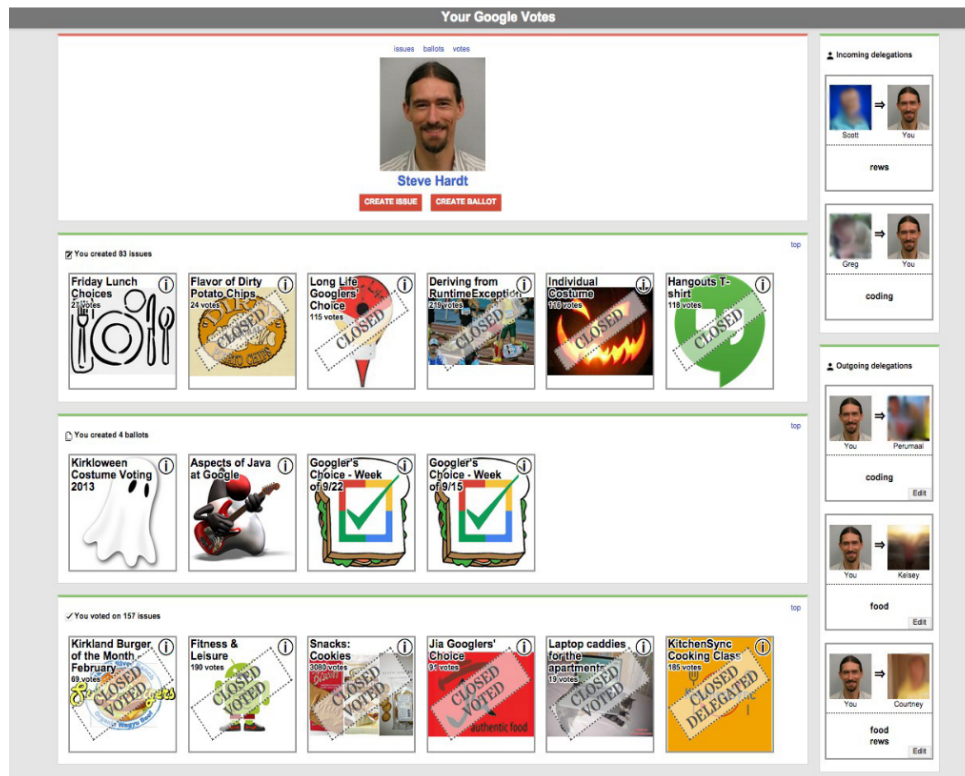
Figure 2.4: Screenshot taken from Hardt and Lopes (2015) showing the user interface of Google Votes.

Google Votes was an internal experiment at Google designed to explore the practical application of liquid democracy within a corporate environment. Built on top of the company's internal Google+ social network, it operated between 2012 and 2015 and allowed employees to participate in decision-making by either voting directly or delegating their vote to a colleague (Hardt and Lopes, 2015).

Delegations in Google Votes were category-specific, meaning that users could choose different delegates for different areas of interest, such as food, events, or technical infrastructure. These delegations were persistent but could be overridden at any time, giving users flexibility to either rely on trusted experts or vote independently as needed. The system supported transitive delegation and allowed users to reclaim control by casting their own vote, even after delegating.

The platform placed strong emphasis on usability and transparency. Delegation features were rolled out incrementally, with additional tools such as voting power estimates and delegation advertisements helping users understand their influence. One key design principle was what the authors called the "Golden Rule of Liquid Democracy": if a user delegates their vote, they should be able to see how it is being used. To accomplish this, users received notifications when their delegate voted, and all

votes were visible to the relevant group. This encouraged accountability and gave voters confidence that their delegated votes were being used appropriately.

While Google Votes was never made publicly available, it served as a successful demonstration of liquid democracy in a structured, real-world setting. It showed that delegative voting could improve engagement and decision-making within large organisations, especially when designed with attention to user experience. For vodle, the system provides a concrete example of how features like topic-specific delegation, transparency tools, and real-time voting feedback can make liquid democracy more practical and accessible.

2.3   Agent Based Modelling

Agent-based modelling (ABM) is a computational approach used to simulate the actions and interactions of autonomous agents in order to assess their effects on a system as a whole. It is particularly suited for exploring complex, dynamic systems where behaviour emerges from local interactions between individual entities (agents) rather than being dictated by central control. ABM has been widely applied in domains such as economics, sociology, and ecology to study decentralised systems, market dynamics, and collective behaviours (Bonabeau, 2002).

The need to explore ABM arises due to the project's goal of introducing a vote-splitting mechanism that hasn't been explored before into vodle. Traditional analysis alone may not effectively capture the dynamic interactions or unintended consequences that can emerge from this novel feature. Through ABM, it is possible to simulate realistic voting scenarios, track delegation chains, identify potential power imbalances, and anticipate challenges. These simulations can reveal performance insights and inform design decisions before implementing the mechanisms within the live platform.

Several widely used ABM frameworks exist, each with their own strengths and drawbacks relevant to this project:

- **NetLogo** (Tisue and Wilensky, 2004) is a highly accessible and widely adopted modelling platform known for its user-friendly graphical interface and ease of learning. It offers rapid prototyping capabilities and excellent visualisation features, allowing clear communication of results. However, very complicated models are not compatible with it.

- **Repast** (Collier, 2003) provides a powerful and versatile suite of tools for building large-scale, computationally intensive simulations. It supports distributed

computing, which is beneficial for extensive delegation networks with potentially thousands of agents. However, Repast has a steep learning curve, which could hinder its compatibility with this heavily time restricted project.

- **Mesa** (Kazil et al., 2020) is an open-source framework written in Python and specifically designed for agent-based modelling. Its advantage lies in its integration with Python's ecosystem of data science libraries. Simulations built with Mesa can easily make use of tools such as NumPy and pandas for efficient data processing, and Matplotlib or Seaborn for visualising model outputs. This compatibility allows for rapid analysis and iteration, while also significantly lowering the learning curve for developers already familiar with Python. Mesa offers a practical balance between usability and computational flexibility, making it well-suited for customisable and moderately large simulations.

- **Agents.jl** (Vahdati, 2019) is a high-performance agent-based modelling framework written in Julia. Due to Julia's speed and efficiency, it is suitable for large-scale and computationally demanding simulations. The framework is designed to be user-friendly, with a syntax that is approachable for those familiar with scientific computing. However, the Julia ecosystem is less mature compared to Python's, which may limit the availability of additional libraries and resources.

Given the time constraints of this project, Mesa offers a practical and efficient solution. Its Python-based interface and straightforward setup allow for rapid development without the overhead of learning a new framework. This ease of use enables more time to be spent designing meaningful experiments and analysing results, rather than configuring tooling.

2.4 Vodle

Vodle is a web-based platform for participatory group decision-making. Users participate in polls that allow them to rate a set of options using sliders. When the poll ends, these ratings are aggregated and the MaxParC rating system is used to determine the final result of the poll.

2.4.1 MaxParC - TODO: rework the summary

MaxParC (Maximum Partial Consensus) is the core rating system used in vodle to aggregate user preferences and determine poll outcomes. Introduced by Heitzig et al. (2024), MaxParC was designed to address common limitations of traditional voting

systems, particularly the tendency for majority rule to overlook minority preferences. Its goal is to balance fairness, consensus, and efficiency in collective decision-making.

In MaxParC, each user rates an option on a scale from 0 to 100. This rating reflects the user's willingness to approve that option based on how many other users also support it. Specifically, a rating of x means that the voter will approve the option if fewer than x% of participants disapprove. A rating of 0 means the option is never approved, while 100 means it is always approved regardless of others' opinions. This structure transforms a simple rating into a conditional approval, allowing for a more nuanced expression of preferences with the potential for compromise.
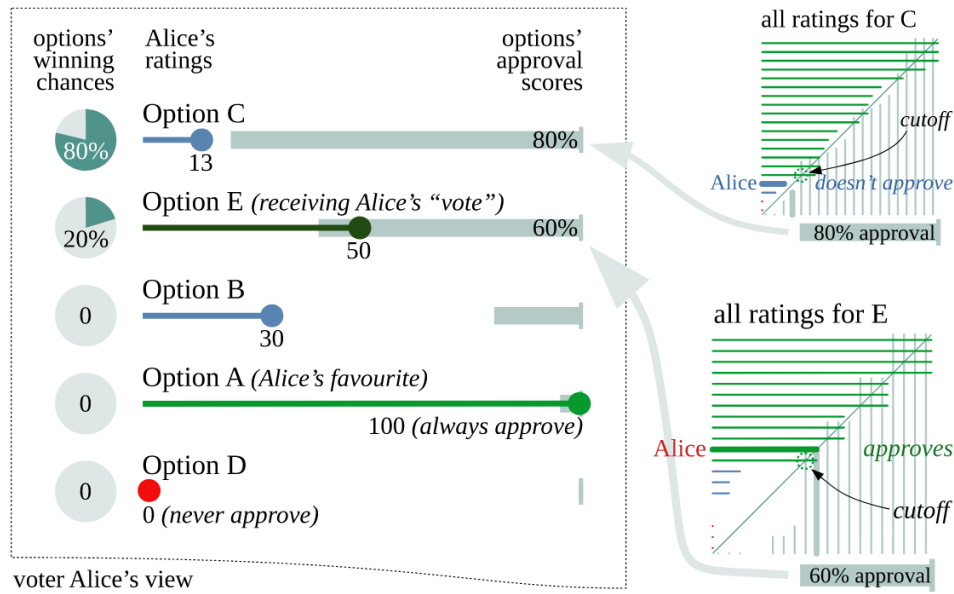
Figure 2.5: Visual representation of MaxParC from the perspective of a voter (Alice). Ratings represent conditional approval thresholds. An option is counted as approved by Alice if the approval bar (light grey) overlaps with her rating needle. Graphic from Heitzig et al. (2024).

TODO: rework this section to be more concise and clear or remove Understanding how MaxParC processes ratings is essential for this project, as the proposed vote splitting mechanism must operate within its conditional approval framework. When a user splits their vote among multiple delegates, the system must ensure that their own rating continues to contribute appropriately. Specifically, if a user delegates x% of their rating to others, their final rating must not fall below $(100 - x)$% of their original input. This constraint guarantees that the user's approval remains proportionally represented, even when part of their voting power is passed on to others.

Integrating liquid democracy into vodle therefore requires careful design to align with MaxParC's logic, ensuring both technical compatibility and conceptual consistency.

2.4.2   Technologies Used

Understanding vodle's technology stack is crucial to successfully integrate liquid democracy features into the existing platform. Since the project involves adding complex delegation and voting logic, it's important to appreciate the constraints and benefits of the technologies currently used in vodle, as they directly influence the design and implementation choices.

Angular

Vodle is built with Angular (Angular Team, 2024), a TypeScript-based frontend framework created by Google. Angular's modularity and structured component system provides a strong foundation for incremental development, essential when introducing new features such as ranked delegation that build upon existing components. Its clear separation of concerns helps maintain readable and maintainable code, simplifying debugging and future enhancements. This is particularly beneficial as the delegation logic is expected to grow in complexity and build upon existing components as the project progresses.

Ionic Framework

The Ionic (Ionic Team, 2024) framework complements Angular by enabling the creation of responsive, mobile-compatible applications from a single codebase. Given vodle's goal of broad user participation, Ionic ensures that its functionalities remain consistent and accessible across both desktop and mobile devices. For this project, new delegation features must be designed to work seamlessly within the existing Ionic framework, ensuring that they are visually appealing and user-friendly on all platforms, including mobile devices.

CouchDB

CouchDB (Apache CouchDB Project, 2024) is vodle's primary data storage method and communicates directly with the client through HTTP requests, meaning there is no dedicated backend. This architecture places significant computational responsibilities on the client-side Angular application, including the handling of delegation chains, cycle detection, and the computation of final vote outcomes. Furthermore, since CouchDB stores data exclusively as JSON-formatted strings, complex delegation structures and voting relationships must be serialised and de-serialised on the client side.

The lack of server-side computation means the delegation algorithms must be designed with client-side efficiency in mind, ensuring performance remains acceptable even as delegation complexity increases. Thus, the choice of algorithms for liquid democracy features, such as those to resolve conflicting delegation paths, is directly influenced by CouchDB's architectural constraints.

These technological considerations shape the practical implementation of liquid democracy within vodle, highlighting the need for efficient client-side processing, careful data management, and cross-platform consistency.

2.4.3 Partially Implemented Delegation in Vodle

The current version of vodle includes a partially implemented delegation feature that allows users to delegate their vote to another participant in a given poll. Whilst this implementation provides the basic structure required for delegation, it lacks consistency in handling more complex delegation scenarios such as cycles.

include delegation workflow + forward ref to implementation chapter

2.4.4 Design Philosophy

Colour pallet - need to maintain consisten design when adding delegation features.

2.5 Summary

The background research presented in this chapter has provided the necessary foundation for designing and implementing advanced delegation features within vodle. Initially, the research highlighted critical limitations in traditional liquid democracy systems, such as the formation of delegation cycles, the risks associated with abstentions, and the disproportionate influence of super-voters. These insights showed the need for implementing delegation mechanisms that are capable of addressing these challenges effectively.

Research into these mechanisms revealed several promising methods. Ranked delegation was found to be an effective approach for reducing the risk of lost votes, with the MinSum delegation rule being particularly suitable due to its clear balance of efficiency, interpretability, and fairness. Vote splitting was identified as a valuable strategy to allow voters greater flexibility by distributing their influence among multiple trusted delegates. Additionally, the concept of delegating different options to distinct delegates was supported by practical experiences from Google Votes, where topic-specific delegations improved user engagement and representation accuracy.

To ensure a thorough evaluation of these delegation mechanisms before full-scale integration, agent-based modelling was selected as a suitable method, with the Mesa framework identified as the most practical choice due to its ease of use, Python compatibility, and flexibility for rapid experimentation.

The technological constraints of vodle itself, especially the reliance on client-side processing due to the CouchDB architecture, demonstrated the need for efficient, lightweight implementation strategies.

These insights collectively define the project's objectives, which are formalised in the following chapter. The objectives are designed explicitly to address the limitations uncovered in the research, ensuring the integration of liquid democracy into vodle is practical, user-friendly, and aligned with established best practices.

Chapter 3

Project Objectives

This chapter outlines the milestones of the project. These objectives were derived from the background research and are designed to address the technical and theoretical challenges identified with traditional delegation systems.

To manage the project effectively, the objectives are divided into two categories: **core objectives**, which form the backbone of the implementation and are essential to meeting the project's main goals, and **extension objectives**, which provide additional insight or value.

Later in the chapter, each objective is broken down into specific functional and non-functional requirements. This structure helps to clarify expectations, guide implementation, and provide clear criteria for evaluating whether each objective has been met.

Core Objectives:

1. **Implement a Core Delegation Model into Vodle:** Build upon the existing, partially implemented delegation code within the vodle platform to create a fully functional system, including resolving key challenges such as cyclic delegations.

2. **Implement Ranked Delegation into Vodle:** Add a backup delegation mechanism to vodle, allowing users to specify up to 3 delegates.

3. **Implement a Vote Splitting Delegation Mechanism into Vodle:** Add functionality to vodle to delegate fractions of their rating to different delegates. Use the *will come back to when research written up* system to calculate final ratings.

4. **Implement the Ability to Delegate Individual Options to Different Users**: Allow users to delegate the ratings of specific options to different delegates.

Extension Objective:

1. **Simulate Delegation Mechanisms:** Perform agent-based modelling to analyse the effectiveness of various delegation systems, including those outlined in Objectives 1, 2 and 3.

3.1 Project Requirements

The project objectives represent high-level goals that must be translated into specific, actionable requirements. This process is essential for clarifying the scope of the work, ensuring comprehensive coverage of each objective, and establishing a structured foundation for both implementation and evaluation.

To support this, requirements are organised into two categories: functional (F) and non-functional (NF). Functional requirements describe the core behaviours and features the system must support, while non-functional requirements define performance, usability, and other quality-related constraints (Sommerville, 2016). Distinguishing between these categories helps ensure that both the system's functionality and overall user experience are properly addressed.

Each requirement is formulated to be measurable and testable. This allows for objective evaluation during development, facilitates verification against the project goals, and helps identify areas for improvement as the system evolves.

3.1.1 Implement a Core Delegation Model into Vodle

Functional Requirements

- **FR1:** The system shall correctly handle the delegation process from invitation to acceptance.

    - **FR1.1:** The system shall allow users to invite others to act as their delegate.

    - **FR1.2:** The system shall allow invited users to accept delegation requests.

    - **FR1.3:** The system shall prevent users from accepting their own delegation invitations.

    - **FR1.4:** The system shall detect and prevent the formation of delegation cycles.

- **FR2:** The system shall provide users with a clear view of their current delegation, including the ability to revoke it at any time.

- **FR3:** The system shall resolve delegations transitively, such that if User A delegates to B and B delegates to C, User C is the final casting voter for A's vote.

- **FR4:** The system shall allow users to override a delegate's decision for specific poll options by submitting a direct vote.

- **FR5:** The user interface for delegation shall be intuitive and accessible, with clear instructions and minimal friction to perform delegation actions.

Non-Functional Requirements

- **NFR1:** All delegation-related data must be stored in a JSON-encoded format, ensuring compatibility with the existing vodle CouchDB database.

- **NFR2:** Any changes to the database schema must be backward compatible, ensuring that existing data is not lost or corrupted during the upgrade process.

- **NFR3:** Any additional data stored in the database must be encrypted, using the same encryption method as the existing data, to ensure user privacy and consistency with the existing system.

- **NFR3:** The system shall preserve user privacy by ensuring that individual voting preferences and delegation choices are not visible to other users. The only information visible to a delegated user shall be the final vote cast on their behalf.

3.1.2   Implement Ranked Delegation into Vodle

Functional Requirements

- **FR1:** The system shall allow users to specify a ranked list of up to 3 delegates for each poll, with the ranking applying to all options within that poll.

- **FR2:** The system shall apply the MinSum delegation rule to resolve each voter's delegation path based on their ranked list of delegates.

- **FR3:** The system shall allow users to override the ranked delegation by submitting a direct vote for specific poll options.

- **FR4:** The system shall provide users with a clear view of their ranked delegation choices, including the ability to alter their rankings or revoke them at any time.

Non-Functional Requirements

- **NFR1:** The system shall ensure that the ranked delegation process does not introduce significant latency in the voting process, maintaining a response time of less than 2 seconds for delegation-related actions when the number of delegates is less than 100.

- **NFR2:** The user interface for ranked delegation shall be intuitive and accessible, with clear instructions and minimal friction to perform delegation actions.

- **NFR3:** The system shall ensure that any data related to ranked delegation is stored in a JSON-encoded format, ensuring compatibility with the existing vodle CouchDB database.

3.1.3 Implement a Vote Splitting Delegation Mechanism into Vodle

Functional Requirements

- **FR1:** The system shall allow users to delegate their vote to multiple delegates simultaneously for a single poll.

- **FR2:** The system shall allow users to assign a weight to each delegate such that the total weight does not exceed 0.99.

- **FR3:** The system shall use the algorithm described in subsection **??** to calculate the final rating for each option based on the weights assigned to each delegate.

- **FR5:** The system shall provide users with a visual interface to edit the weights assigned to each delegate, with either sliders or numeric inputs for easy adjustment.

Non-Functional Requirements

- **NFR1:** The system shall ensure that vote-splitting calculations are performed entirely on the client side to comply with vodle's CouchDB architecture.

- **NFR2:** All related data must be serialised as JSON strings to ensure compatibility with the CouchDB backend.

- **NFR3:** The user interface for vote splitting shall be intuitive and allow users to adjust weights easily, using sliders or numeric inputs.

3.1.4 Implement the Ability to Delegate Individual Options to Different Users

Functional Requirements

- **FR1:** The system shall allow users to assign different delegates for each individual option or subset of options within a poll.

- **FR2:** The system shall ensure that each delegated option is resolved independently, using the appropriate delegate's vote for that option.

- **FR3:** The system shall allow users to override a delegate's vote for a specific option by submitting their own rating.

- **FR4:** The system shall provide a user interface for viewing or revoking each individual delegation.

Non-Functional Requirements

- **NFR1:** The delegation interface must be intuitive and clearly indicate which delegate is assigned to each option, ensuring ease of use.

- **NFR2:** The delegation data must be serialised in a format compatible with CouchDB (e.g., JSON-encoded) to maintain compatibility with vodle's storage system.

3.1.5 Simulate Delegation Mechanisms

Functional Requirements

- **FR1:** The simulation system shall model individual agents representing voters, each capable of voting, abstaining, or delegating their vote according to a selected delegation rule.

- **FR2:** The system shall support multiple delegation mechanisms, including standard transitive delegation, ranked delegation (with the MinSum delegation rule), and vote splitting.

- **FR3:** The system shall allow configuration of simulation parameters such as number of agents, delegation probabilities and abstention rates.

- **FR4:** The system shall track and record key metrics such as vote concentration, number of super-voters, average chain length, vote loss due to abstentions or cycles, and decision quality.

- **FR5:** The system shall output simulation results in a structured format (e.g. CSV or JSON) for further analysis.

Non-Functional Requirements

- **NFR1:** The simulation framework must be lightweight and easy to extend, enabling rapid experimentation with new delegation rules or metrics.

- **NFR2:** The system shall be developed using Mesa to take advantage of existing data science libraries such as NumPy, Pandas, and Matplotlib for analysis and visualisation.

- **NFR4:** The simulation design shall support reproducibility by enabling fixed random seeds and storing configuration settings alongside output data.

Chapter 4

Design and Implementation

This chapter outlines the design choices and technical implementation of the delegation mechanisms introduced into the vodle platform. Each section corresponds to a key objective defined in the previous chapter, detailing how it was approached within the constraints of vodle's architecture and data model.

4.1 System Architecture Overview

Vodle is built as a serverless web application that emphasises accessibility, client-side performance, and ease of deployment. Its architecture comprises two components:

- **Frontend:** Implemented using Angular and the Ionic framework, the frontend provides a responsive and modular interface that works across both desktop and mobile devices. The use of Angular facilitates the creation of component-based user interfaces, essential for introducing interactive features such as the ranked delegation UI and vote splitting sliders.

- **Backend:** Vodle uses CouchDB as its database. There is no custom backend logic or middleware; instead, the application communicates directly with CouchDB over HTTP. All documents, including user ratings and delegation structures, are stored in JSON format.

- **Implications of a serverless architecture:**

  - All vote delegation logic, including transitive resolution, cycle detection, and vote splitting calculations, must be executed in the browser. This places constraints on performance and requires careful optimisation of algorithms used.

– The lack of a traditional server-side backend limits opportunities for centralised verification or caching, increasing the complexity of handling dynamic delegation graphs client-side.

– Data structures must be serialised and de-serialised as JSON, requiring careful schema design to ensure backwards compatibility and clarity across different types of delegations.

4.1.1 CouchDB overview

4.2 Implement a Core Delegation Model into vodle

This section provides a self-contained, technical narrative of how the core delegation mechanism was conceived, which problems it solved, and how the final client-side implementation works. It follows the structure below so that each stage (from user interaction to data consistency) can be expanded or shortened independently during editing.

4.2.1 Delegation Interaction Flow

Figure 4.1 depicts the life-cycle of a delegation. The diagram is intentionally simple–only the four message exchanges required to set up or tear down a delegation are shown–so the reader can map each arrow to a concrete API call or UI control.

accept

Delegator A share link Delegate B

tally uses B   casts vote

Figure 4.1: Minimal message sequence for creating a delegation. Only four interactions are required: (1) A shares a link, (2) B accepts, (3) B may cast or delegate further, (4) A's vote resolves through B.

At runtime the app must therefore support:

1. Generating and copying a poll-scoped, user-scoped invitation link.

2. Polling CouchDB for acceptance (or rejection) of that link.

3. A client-side resolution routine that follows the chain of `delegateFor` relationships until a casting voter is found.

4. Revocation that is immediate yet conflict-free if multiple clients are open.

4.2.2 Current Issues Prior to Redesign

Out-of-sync state Each browser session cached its own copy of the delegation graph; no push notifications existed to bring other clients up-to-date.

Inconsistent cycle checking Because state diverged, the local "no-loop" tests sometimes produced different answers on different machines.

UI drift Controls such as "delegate none/some/all" and the "your vote is used for n others" badge used stale data and became misleading.

4.2.3 Cycle Checking

The following section details the design and implementation of a cycle-checking algorithm that was added to the core delegation model as well as the data structures and synchronisation mechanisms that were introduced to support it.

Algorithm Design

The current delegations in a system can be represented as a directed graph where each user is represented as a node and each delegation is represented as a directed edge (u, v), where u is the delegator and v is the delegate. The goal of the cycle-checking algorithm is to ensure that a proposed delegation does not create a cycle in this directed graph.

A new delegation $X \rightarrow Y$ is valid if and only if Y is *not* reachable from X in that DAG – if Y is not a descendant of X.

Instead of checking for this condition directly using a depth-first search (DFS) or breadth-first search (BFS), a more efficient approach is to maintain a list of all descendants for each user. This allows us to check if Y is in the list of descendants of X in constant time. This implementation will be detailed in the following section.

Implementation Details

A hashmap is used to store the descendants of each user. The keys are user IDs, and the values are sets of user IDs representing the direct delegates of that user. In the code, this hashmap is referred to as "`inverse_indirect_map`".

```
"inverse_indirect_map": {
  "B": ["A"],
  "C": ["B", "A"],
  "D": ["C", "B", "A"]
}
```

Figure 4.2: Example of a hashmap for users A, B, C, and D. User A has delegated to B, user B has delegated to C, and user C has delegated to D. Consequently, the descendants of user D are A, B and C.

This structure allows for the following operations:

Check Delegation   To check if a delegation $X \to Y$ is valid, we simply check if Y is in the set of descendants of X. This operation takes $O(1)$ time.

```
const inverse_indirect_map = this.G.D.get_inverse_indirect_map(pid);
const descendant_set = inverse_indirect_map.get(delegate_vid);
if (descendant_set.has(myvid)) {
  cycle = true;
}
```

Figure 4.3: Code for checking if a delegation is valid. This check is done when a user clicks on a delegate link. The map is retrieved from the synced local cache, and the set of descendants for the delegator is checked to see if it contains the current user.

- **Check Delegation**: to decide if $X \to Y$ is legal, verify that `Y` $\notin$ `map[X]` ($O(1)$).

- **Add Edge**: `map[Y]` $\leftarrow$ `map[Y]` $\cup$ `map[X]` $\cup$ `{X}`.

- **Remove Edge**: `map[Y]` $\leftarrow$ `map[Y]` $\setminus$ (`map[X]` $\cup$ `{X}`). Because the graph is a DAG, it is not possible for the descendants of X to be a subset of the descendants of Y after removing the edge $X \to Y$.

Persistence   The map lives alongside the poll document so every client shares one authoritative copy. Although ordinary poll docs are append-only, a special admin rule now allows updates at the key `inverse_indirect_map`. Conflict fields (`_rev` in CouchDB) ensure last-write-wins semantics while we serialise edits.

Practical Challenges

- blank 4 now.

4.2.4   UI Fixes

- Re-bound the "delegate none/some/all" switch so it toggles *only* the current option and refreshes its label after server confirmation.

- Re-implemented the "vote is used for *n* others" badge so it recomputes from the latest `inverse_indirect_map` rather than client guesses.

- Added unobtrusive toast notifications when a delegation request is rejected (e.g. cycle detected) so the user can immediately choose another delegate.

4.2.5   Summary

The redesign introduced a lightweight yet provably correct cycle-checking algorithm, synchronised across clients by persisting an `inverse_indirect_map` in the poll document. Combined with targeted UI fixes, the core delegation model now (i) prevents loops, (ii) reflects state consistently across devices, and (iii) scales to thousands of voters without central servers.

4.3   Implement Ranked Delegation into Vodle

This feature introduced ranked delegation using the MinSum rule, allowing users to list fallback delegates in case their primary choice was unavailable.

- UI for setting delegate rankings

- `direct_delegation_map`: maps user IDs to list of ranked delegates [[delegationid, rank, status]...]

- Explanation and application of the MinSum rule

- Implementation of ranked path resolution

- Illustrations and code snippets

Challenges

The MinSum rule had to be implemented efficiently using only browser-based resources. Ranking resolution had to preserve user intent while avoiding delegation ambiguity. Providing visual feedback to help users understand how rankings would resolve added an additional layer of design complexity.

4.4 Implement a Vote Splitting Delegation Mechanism into Vodle

Vote splitting was implemented to allow users to distribute fractional influence to multiple delegates.

- UI for assigning weights

- modify`direct_delegation_map`to include weights [[delegationid, weight, status]...]

- Computation of weighted vote outcomes

- Constraints:

  - Weight sum limit (`< 1.0`)
  - Error handling

- Algorithm integration and frontend testing

Challenges

The vote splitting logic needed to maintain consistency with the MaxParC aggregation model, while ensuring intuitive user experience. Edge cases (e.g., partially overlapping delegate chains or missing data) introduced complexity during testing. Rendering weight distributions clearly in the UI while keeping the interface lightweight was a recurring challenge.

4.5 Implement the Ability to Delegate Individual Options to Different Users

This feature enabled per-option delegation, allowing users to assign a different delegate for each item in a poll.

- Per-option delegate selection interface

- Independent resolution of each delegated option

- talk about nested map - need to take care to serialise.

- `direct_delegation_map: option_id -> user_id -> [delegationid, null, status]`

- `inverse_indirect_map: option_id -> user_id ->` list of users who have delegated to them, either directly or indirectly.

- Storage schema modifications

Challenges

This mechanism required updates to the internal delegation logic to handle resolution at the option level. The user interface also had to be adapted to display multiple concurrent delegate selections without overwhelming the user. Debugging resolution logic for hybrid delegation modes (e.g., one direct, one split, one ranked) was non-trivial.

4.6   Simulate Delegation Mechanisms

The simulation objective was de-scoped due to time constraints and prioritisation of implementation work. While initial planning and framework selection (Mesa) were completed, no functional simulation code was delivered. The decision to drop this extension is discussed further in the Project Management chapter.

4.7   Design Decisions and Trade-offs

- All logic had to run client-side due to the serverless CouchDB architecture, limiting complexity and computational resources.

- A consistent JSON format was required for all data models, impacting flexibility in data design.

- Trade-offs were made between expressive delegation types and usability, particularly in the option-specific and vote splitting interfaces.

4.8   Summary

- Each objective was successfully implemented within the constraints of the vodle platform.

- Challenges were primarily technical (client-side performance, real-time resolution) and design-oriented (clarity and control for users).

- The final implementation offers a modular, extensible delegation system that addresses the key theoretical and practical limitations outlined in earlier chapters.

Chapter 5

Evaluation

5.1 Testing

- unit testing for delegation algorithms (minsum, vote splitting and test cycle checking for standard delegation)

-

- make sure all requirements are done (can format as table)

5.2 Feedback From Customer

get quote from Jobst: what he likes and dislikes about the implementation. also include specifications about what the final delegation implementation will be (vote splitting).

5.3 Self Reflection

"It can be a useful exercise for you (and a point of consolidation for the reader) to put together a brief summary of what you have achieved. This is not a compulsory section, but a self-assessment is welcome. A suggested format for this is to include a short section entitled 'Author's Assessment of the Project' consisting of brief (up to 100 words) answers to each of the following questions.

- What is the (technical) contribution of this project?

- Why should this contribution be considered relevant and important for the subject of your degree?

- How can others make use of the work in this project?

- Why should this project be considered an achievement?

- What are the limitations of this project?

"

Chapter 6

Project Management

This chapter outlines the project's management approach, including the development methodology, planning, and reflections on the process. It also considers legal and ethical issues and assesses key risks associated with the project.

6.1   Methodology

The project followed an agile methodology, selected for its flexibility and its emphasis on iterative development and regular customer feedback. The work involved building a sequence of interdependent features into vodle; starting with a basic delegation mechanism and expanding to include ranked delegation, vote splitting, and per-option delegation. Since each feature built on the last, an iterative approach was ideal to ensure compatibility and to allow design decisions to adapt over time.

Agile was particularly well suited to this project due to the presence of an active "customer" figure: Jobst Heitzig, the co-supervisor and original creator of vodle. Heitzig provided clarified system expectations and helped shape design decisions based on the real-world use case. Fortnightly meetings were held with both Jobst Heitzig and Markus Brill, to review progress and incorporate their feedback into the next development cycle. This tight feedback loop is a core principle of agile, allowing the project to stay aligned with user needs and system goals.

Alternative models such as Waterfall were considered but ultimately rejected due to their inflexibility. Waterfall could have provided clear documentation at each phase and well-defined milestones, which initially seemed beneficial for implementing interconnected voting delegation features. However, Waterfall requires defining the full scope of the project upfront and offers limited room for revision - something that would have been impractical due to the project's shorter timeframe and the evolving nature of requirements as features were tested.

Elements of Scrum were also considered, as it is one of the most widely adopted agile frameworks (63% of agile teams use Scrum (VersionOne, 2020)). The sprint-based development cycles, clearly defined roles, and structured ceremonies like sprint planning and reviews were attractive features that could have facilitated focused implementation and effective communication. However, daily stand-up meetings and fixed-length sprints were not feasible for this project, as all parties involved (myself and the two supervisors) had other commitments. The small team size didn't warrant all Scrum ceremonies, and the academic nature of the project called for more flexible review cycles. Instead, progress was reviewed every two weeks, ensuring feedback could still be gathered and acted upon without adding unnecessary scheduling pressure.

move to start? Each development cycle produced a working, testable feature that could be evaluated and integrated into the overall system. This approach reduced the risk of late-stage errors and helped maintain steady progress throughout the project. Agile's iterative and feedback-driven structure was a natural fit for the technical and collaborative demands of this work.

6.2   Plan

The project plan was organised into objectives (see section 3) that built on one another in sequence:

- **Core Objective 1:** Implement a Core Delegation Model into Vodle.

- **Core Objective 2:** Implement Ranked Delegation into Vodle.

- **Core Objective 3:** Implement a Vote Splitting Delegation Mechanism into Vodle.

- **Core Objective 4:** Implement the Ability to Delegate Individual Options to Different Users.

- **Extension Objective 1:** Simulate Delegation Mechanisms.

This objective-led structure was well-suited to the agile approach, allowing each milestone to be treated as an iteration with a deliverable at the end. A Gantt chart (see below) was created to visualise the project timeline and to track dependencies and progress.

TODO: Gantt Chart

6.3   Risk Assessment - TODO

Risk	Likelihood	Mitigation Strategy
Breaking the live vodle site during development	Medium	Use Git branching to isolate development from production environments. Conduct local testing before deployment.
Feature complexity exceeds estimates	High	Prioritise core objectives and maintain flexibility in scope.
Lack of engagement from supervisors or stakeholders	Low	Maintain regular communication through scheduled meetings.
Data loss or corruption	Low	Use Git for version control and take regular local backups.

Table 6.1: Key risks identified and their mitigation strategies

The following sections provide a detailed analysis of the risks identified in the risk assessment table (see Table 6.1). Each risk is discussed individually, outlining its implications and the strategies proposed to mitigate potential issues.

Breaking the Live Vodle Site During Development

Likelihood: Medium

Description: Modifications to the existing vodle platform could unintentionally introduce downtime or impair existing functionality on the live website. Any disruptions could negatively impact real users' interactions, leading to dissatisfaction and loss of trust in the platform.

Mitigation: Development activities will utilise Git branching to isolate new code from the production environment. Features will be developed and rigorously tested in local or staging environments before integration with the live deployment. Incremental rollouts and thorough pre-deployment testing will further help identify potential problems early, allowing quick remediation or rollback.

Feature Complexity Exceeds Estimates

Likelihood: High

Description: Advanced features such as ranked delegation and vote splitting may prove more complex than initially anticipated. Unexpected complexity can lead to delays, reduced functionality, or incomplete implementations, potentially affecting the project's timeline and deliverables.

Mitigation: Core objectives have been clearly defined and prioritised, ensuring focus remains on essential functionality. In cases of higher than anticipated complexity, resources will be redirected towards completing critical core features first, while the extension objective (agent based modelling) can be scaled back or postponed as needed. Regular agile reviews will monitor progress closely, facilitating early identification and management of complexity-related issues.

Lack of Engagement from Supervisors or Stakeholders

Likelihood: Low

Description: Regular feedback and engagement from supervisors and stakeholders are crucial to ensure alignment with project goals, requirements, and user expectations. Insufficient feedback could result in misaligned implementations or objectives that do not fully meet user needs.

Mitigation: Fortnightly meetings have been scheduled with both the primary supervisor (Markus Brill) and the co-supervisor (Jobst Heitzig), who also fulfils the role of the project customer. This structured schedule ensures consistent opportunities for input and feedback. Additionally, a Telegram group chat is available to handle urgent queries and maintain ongoing dialogue.

Data Loss or Corruption – Code

Likelihood: Low

Description: Development activities pose a risk of code loss or corruption due to accidental deletion, unintended changes, or version conflicts. Such incidents could significantly delay development and necessitate additional time for recovery.

Mitigation: Version control will be rigorously maintained using Git, with frequent commits and descriptive commit messages ensuring traceability. Regular backups of the repository will be taken to safeguard against accidental loss, providing straightforward recovery paths when needed.

Data Loss or Corruption – Database

Likelihood: Low

Description: While unlikely, corruption of the CouchDB database during development could occur due to improper schema modifications or accidental changes.

Mitigation: No mitigation — in the event of data corruption, the development database will be reset to its original state. As all data in the database is poll and user

specific, it does not impact development as no important data is stored in the production environment.

6.4   Legal and Ethical Considerations

As vodle may eventually be used to gather votes on sensitive topics, care was taken to ensure privacy and fairness. Delegation chains are resolved internally and never publicly exposed, preserving the confidentiality of voter relationships. No personal data was collected or processed for the purposes of this project, so no changes to the platform's terms of service were required.

6.5   Risk Management Reflection - TODO

6.6   Overall/Self Reflection - TODO

Chapter 7

Conclusions

7.1 Future Work

Bibliography

Angular Team, . *Angular - Web Framework*. Google, 2024. URL `https://angular.io/`. `https://angular.io/`.

Apache CouchDB Project, . *Apache CouchDB*. The Apache Software Foundation, 2024. URL `https://couchdb.apache.org/`. Accessed April 2025. `https://couchdb.apache.org/`.

Behrens, Jan & Swierczek, Björn. Preferential delegation and the problem of negative voting weight. *The Liquid Democracy Journal*, 3:6–34, 2015.

Behrens, Jan & Kistner, Axel & Nitsche, Andreas & Swierczek, Björn. *The Principles of LiquidFeedback*. Interacktive Demokratie, 2014.

Bersetche, Francisco M. Generalizing Liquid Democracy to multi-agent delegation: A Voting Power Measure and Equilibrium Analysis, April 2024.

Blum, Christian & Zuber, Christina Isabel. Liquid Democracy: Potentials, Problems, and Perspectives. *Journal of Political Philosophy*, 24(2):162–182, June 2016. ISSN 0963-8016, 1467-9760. doi: 10.1111/jopp.12065. URL `https://onlinelibrary.wiley.com/doi/10.1111/jopp.12065`.

Bonabeau, Eric. Agent-based modeling: Methods and techniques for simulating human systems. *Proceedings of the National Academy of Sciences*, 99:7280–7287, 2002.

Brill, Markus & Delemazure, Théo & George, Anne-Marie & Lackner, Martin & Schmidt-Kraepelin, Ulrike. Liquid Democracy with Ranked Delegations. *Proceedings of the AAAI Conference on Artificial Intelligence*, 36(5):4884–4891, June 2022. ISSN 2374-3468. doi: 10.1609/aaai.v36i5.20417. URL `https://ojs.aaai.org/index.php/AAAI/article/view/20417`. Number: 5.

Christoff, Zoé & Grossi, Davide. Liquid democracy: An analysis in binary aggregation and diffusion, 2017. URL `https://arxiv.org/abs/1612.08048`.

Collier, Nick. Repast: An extensible framework for agent simulation. *The University of Chicago's social science research*, 36:2003, 2003.

Degrave, Jonas. Resolving multi-proxy transitive vote delegation. (arXiv:1412.4039), December 2014. doi: 10.48550/arXiv.1412.4039.

Ford, Bryan Alexander. Delegative Democracy. May 2002. URL `https://infoscience.epfl.ch/handle/20.500.14299/156450`.

Hall, Andrew B. & Miyazaki, Sho. What Happens When Anyone Can Be Your Representative? Studying the Use of Liquid Democracy for High-Stakes Decisions in Online Platforms. Technical report, 2024.

Hardt, Steve & Lopes, Lia. Google Votes: A Liquid Democracy Experiment on a Corporate Social Network. *Defensive Publications Series*, June 2015.

Heitzig, Jobst & Simmons, Forest W. & Constantino, Sara M. Fair group decisions via non-deterministic proportional consensus. *Social Choice and Welfare*, May 2024. ISSN 1432-217X. doi: 10.1007/s00355-024-01524-3. URL `http://dx.doi.org/10.1007/s00355-024-01524-3`. Publisher: Springer Science and Business Media LLC.

Ionic Team, . *Ionic Framework - Cross-Platform Mobile App Development*. Ionic, 2024. URL `https://ionicframework.com/`. Accessed April 2025. `https://ionicframework.com/`.

Kazil, Jackie & Masad, David & Crooks, Andrew. Utilizing Python for Agent-Based Modeling: The Mesa Framework. In Thomson, Robert & Bisgin, Halil & Dancy, Christopher & Hyder, Ayaz & Hussain, Muhammad, editors, *Social, Cultural, and Behavioral Modeling*, pages 308–317, Cham, 2020. Springer International Publishing. ISBN 978-3-030-61255-9.

Kling, Christoph Carl & Kunegis, Jerome & Hartmann, Heinrich & Strohmaier, Markus & Staab, Steffen. Voting behaviour and power in online democracy: A study of liquidfeedback in germany's pirate party, 2015. URL `https://arxiv.org/abs/1503.07723`.

Kotsialou, Grammateia & Riley, Luke. Incentivising Participation in Liquid Democracy with Breadth-First Delegation. *New Zealand*, 2020.

Sommerville, Ian. *Software Engineering*. Always Learning. Pearson, Boston Columbus Indianapolis New York San Francisco Hoboken Amsterdam Cape Town Dubai London, tenth edition edition, 2016. ISBN 978-1-292-09613-1 978-1-292-09614-8.

Sven Becker, DER SPIEGEL. Liquid Democracy: Web Platform Makes Professor Most Powerful Pirate - spiegel.de, 2012. URL `https://www.spiegel.de/international/germany/liquid-democracy-web-platform-makes-professor-most-powerful-pirate-a-818683.html`.

Tisue, Seth & Wilensky, Uri. Netlogo: A simple environment for modeling complexity. In *International Conference on Complex Systems*, volume 21, pages 16–21. Citeseer, 2004.

Vahdati, Ali. Agents.jl: Agent-based modeling framework in Julia. *Journal of Open Source Software*, 4(42):1611, October 2019. ISSN 2475-9066. doi: 10.21105/joss.01611.

VersionOne, . 14th annual state of agile report. `https://stateofagile.com`, 2020. Accessed: 2025-04-09.